

HPE DSI 311


Introduction to Machine Learning

Summer 2021

Instructor: Ioannis Konstantinidis

Overview

- Simple models to get started:
 - K nearest neighbors
 - Logistic Regression
- Hands-on examples
 - Jupyter Notebooks
 - Data Summarization



ML
techniques:
How do they
differ?



ML techniques by data type

Unsupervised: aims to uncover groups of observations from initially unclassified data

Analyze: How is the data set X structured?

ML techniques by data type

Unsupervised: aims to uncover groups of observations from initially unclassified data

Analyze: How is the data set X structured?

Supervised: works with data that is already classified to tailor rules for classifying new (and as yet unclassified) individuals; no feedback

Predict: What would the data point x do?

ML techniques by data type

Unsupervised: aims to uncover groups of observations from initially unclassified data

Analyze: How is the data set X structured?

Supervised: works with data that is already classified to tailor rules for classifying new (and as yet unclassified) individuals; no feedback

Predict: What would the data point x do?

Reinforcement learning: the problem faced by an agent that learns behavior through trial-and-error interactions with a dynamic environment

Gameify: Faced with (X,Y) , which strategy wins?

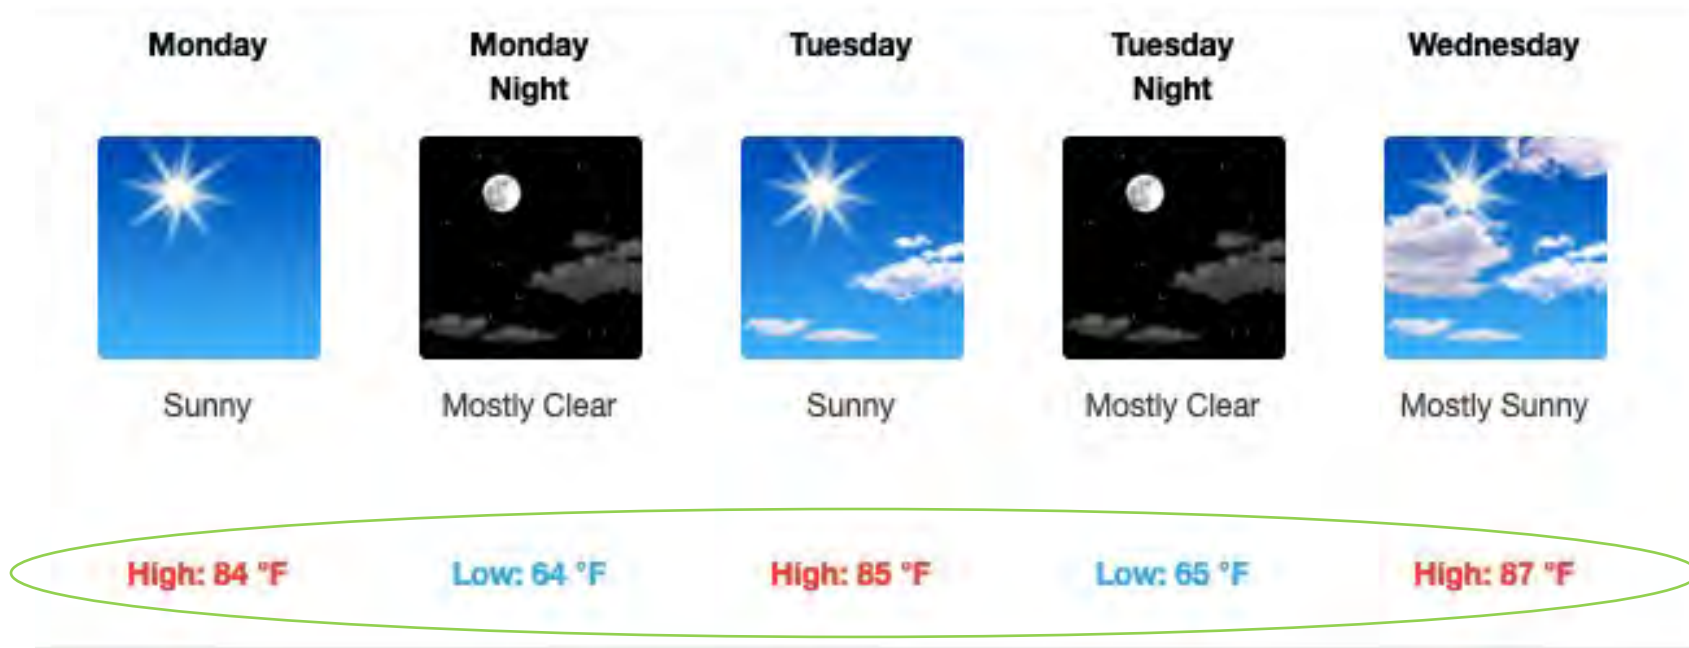
ML techniques by data type

Unsupervised	Supervised	Reinforcement Learning
Clustering / Anomaly detection	Classification/ Regression	Dynamic Programming (Monte Carlo, genetic algorithms)
Signal Separation (matrix factorization)		Deep RL

What types
of
classification
problems?

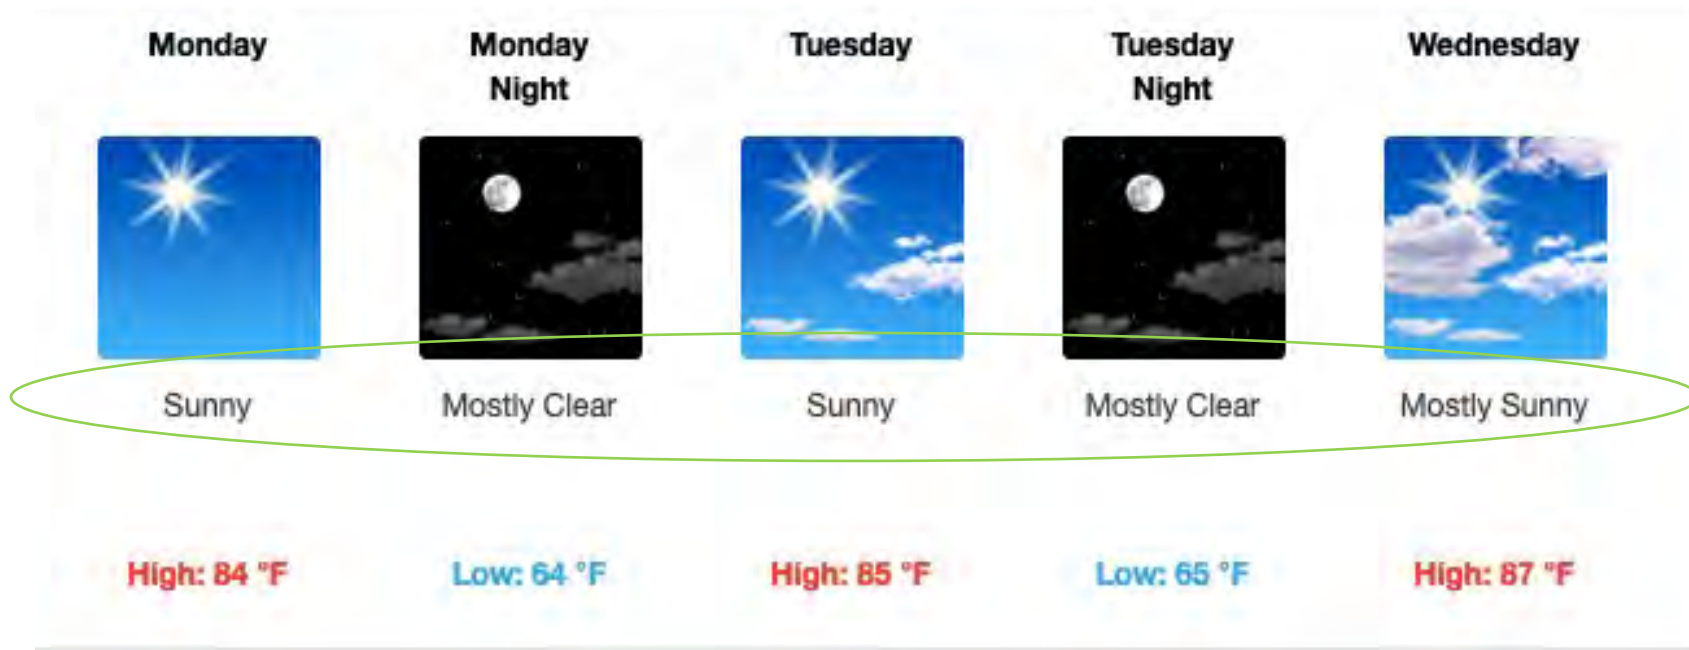


Supervised ML



Regression:
predict a number

Supervised ML



Classification:
predict a label

Regression:
predict a number

Classification using K-Nearest Neighbors

accepting (word
article).

focus n point

converging rays of light,
heat, waves of sound meet;

center of activity or
intensity; foci focal
adjust; cause to converge;
concentrate; a focal
pertaining to focus

K-Nearest Neighbors: algorithm

Training algorithm

- All training example points (x_{train}, y_{train}) go into a reference list

Classification algorithm

- Given a query instance x_{test} to be classified, find the **nearest** point x_{train} in the reference list
- Repeat until **k nearest points** are identified from the list
- Calculate $y_{predict}$ based on the values of y_{train} for these neighbors, i.e., k nearest points


In practical terms: training

	Known input			Known labels

In practical terms: prediction

	Nearest neighbor			predicted label
	New input			

x



Hands-on
Example:

k-NN

K-Nearest Neighbors: sklearn implementation

```
KNeighborsClassifier(n_neighbors=5,  
weights='uniform',  
algorithm='auto',  
leaf_size=30,  
p=2,  
metric='minkowski',  
metric_params=None,  
n_jobs=None,  
**kwargs)
```

<https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>

K-Nearest Neighbors: choice of **K**


```
KNeighborsClassifier(n_neighbors=5,  
weights='uniform',  
algorithm='auto',  
leaf_size=30,  
p=2,  
metric='minkowski',  
metric_params=None,  
n_jobs=None,  
**kwargs)
```

K-Nearest Neighbors: choice of K

Who are the neighbors of the new sample (green circle)?

Blue squares or red triangles?

$$g(\mathbf{x}) = \sum_{i \in k\text{NN}(\mathbf{x})} y_i$$



- $k = 1$: a RED TRIANGLE is the nearest neighbor, so the guess would be RED TRIANGLE
- $k = 3$ (solid line circle): 2 red triangles and only 1 blue square in the neighborhood, so the guess would be RED TRIANGLE
- $k = 5$ (dotted line circle): 3 blue squares and only 2 red triangles in the neighborhood, so the guess would be BLUE SQUARE

K-Nearest Neighbors: choice of **weight**

```
KNeighborsClassifier(n_neighbors=5,  
weights='uniform',  
algorithm='auto',  
leaf_size=30,  
p=2,  
metric='minkowski',  
metric_params=None,  
n_jobs=None,  
**kwargs)
```

K-Nearest Neighbors: choice of **weight**

Weights default='uniform'

weight function used in prediction. Possible values:

- 'uniform' : uniform weights. All points in each neighborhood are weighted equally.
- 'distance' : weight points by the inverse of their distance. In this case, closer neighbors of a query point will have a greater influence than neighbors which are further away.
- [callable] : a user-defined function which accepts an array of distances, and returns an array of the same shape containing the weights.


K-Nearest Neighbors: choice of **weight**

Who are the neighbors of the new sample (**green circle**)?

Blue squares or **red triangles**?

$$g(\mathbf{x}) = \sum_{i \in \text{kNN}(\mathbf{x})} \text{weight}(\mathbf{x}_i, \mathbf{x}) y_i$$

Default option for $\text{weight}(\mathbf{x}_i, \mathbf{x})$ is uniform (every weight = 1)




K-Nearest Neighbors: choice of **weight**

Who are the neighbors of the new sample (**green circle**)?

Blue squares or **red triangles**?

$$g(\mathbf{x}) = \sum_{i \in \text{kNN}(\mathbf{x})} \text{weight}(\mathbf{x}_i, \mathbf{x}) y_i$$



k = 5:

3 distant **blue squares** and 2 close **red triangles** in the neighborhood

- Uniform weights: the guess would be **BLUE SQUARE**
- Distance weights: the guess would be **RED TRIANGLE**


K-Nearest Neighbors: choice of **metric**

```
KNeighborsClassifier(n_neighbors=5,  
weights='uniform',  
algorithm='auto',  
leaf_size=30,  
p=2,  
metric='minkowski',  
metric_params=None,  
n_jobs=None,  
**kwargs)
```

K-Nearest Neighbors: choice of **metric**

Pythagorean theorem:

$$\begin{aligned} & \{ \text{distance}(\mathbf{X}_1, \mathbf{X}_2) \}^2 \\ &= \\ & \{ \text{distance_along_u_axis}(\mathbf{X}_1, \mathbf{X}_2) \}^2 \\ &+ \\ & \{ \text{distance_along_v_axis}(\mathbf{X}_1, \mathbf{X}_2) \}^2 \end{aligned}$$



$$\{\text{distance}(\mathbf{X}_1, \mathbf{X}_2)\}^2 = (u_1 - u_2)^2 + (v_1 - v_2)^2$$

Euclidean: As the crow flies

Dimension(X) = 2


Example:

V axis	2	2.236	2.828
	1	1.414	2.236
	X_1	1	2
U axis			

Manhattan: As a Lyft drives

$$\begin{aligned} \text{distance}(\mathbf{X}_1, \mathbf{X}_2) \\ = \\ \text{distance_along_u_axis}(\mathbf{X}_1, \mathbf{X}_2) \\ + \\ \text{distance_along_v_axis}(\mathbf{X}_1, \mathbf{X}_2) \end{aligned}$$

Dimension(X) = 2



$$\text{distance}(\mathbf{X}_1, \mathbf{X}_2) = |\mathbf{u}_1 - \mathbf{u}_2| + |\mathbf{v}_1 - \mathbf{v}_2|$$

Maximum Distance


$$\begin{aligned} &\text{distance}(\mathbf{X}_1, \mathbf{X}_2) \\ &= \\ &\text{Max}\{ \text{distance_along_u_axis}(\mathbf{X}_1, \mathbf{X}_2), \\ &\quad \text{distance_along_v_axis}(\mathbf{X}_1, \mathbf{X}_2) \} \end{aligned}$$

Dimension(X) = 2

V axis	2	2	2
	1	1	2
	\mathbf{X}_1	1	2
U axis			

Minkowski: The L^p metric

$$\begin{aligned} & \{ \text{distance}(\mathbf{X}_1 , \mathbf{X}_2) \}^p \\ &= \\ & \{ \text{distance_along_u_axis}(\mathbf{X}_1 , \mathbf{X}_2) \}^p \\ &+ \\ & \{ \text{distance_along_v_axis}(\mathbf{X}_1 , \mathbf{X}_2) \}^p \end{aligned}$$



$p = 2$: Euclidean – Each point on blue arc is same distance from LL corner

$p = 1$: Manhattan – Each point on violet diagonal is same distance from LL corner

$p = \infty$: Maximum – Each point on red sides is same distance from LL corner

Minkowski: The L^p metric

Metric is the choice of distance for finding the nearest neighbors. The default metric is minkowski with $p=2$, which is equivalent to the standard Euclidean metric.

P Power parameter for the Minkowski metric.

- When $p = 1$, this is equivalent to using `manhattan_distance` (l_1)
- When $p=2$, this is `euclidean_distance` (l_2)
- For arbitrary p , it is `minkowski_distance` (l_p)

K-Nearest Neighbors: choice of **algorithm**

```
KNeighborsClassifier(n_neighbors=5,  
weights='uniform',  
algorithm='auto',  
leaf_size=30,  
p=2,  
metric='minkowski',  
metric_params=None,  
n_jobs=None,  
**kwargs)
```

K-Nearest Neighbors: choice of **algorithm**

Algorithm default='auto'

Algorithm used to find the nearest neighbors:

- 'ball_tree' will use a [BallTree](#) algorithm
- 'kd_tree' will use a [KDTree](#) algorithm
- 'brute' will use a brute-force search
- 'auto' will attempt to decide the most appropriate algorithm based on the values passed to [fit](#) method

K-Nearest Neighbors: AKA lazy, instance-based learning

Lazy: No training process

Instance-based: Construct only local approximation to the target function that differs based on the neighborhood of each new query instance

Are there any disadvantages?

Cost of classifying new instances can be high:

- Nearly all computation takes place at classification time rather than learning time
- Number of points needed for good coverage of feature space scales exponentially with number of dimensions

What are
some other
ML
methods?



Logistic Regression

accepting (word
article).

focus n point

converging rays of light,

heat, waves of sound, meet;

centre of activity or

intensity; pl focuses, foci; v

adjust; cause to converge;

concentration; focal

pertaining to focus


VERY simple example

HOURS STUDYING	PASSED EXAM
4	YES
1	NO
3.5	NO
2.25	YES
0.25	NO

Known
input

Known
labels

VERY simple example




VERY simple example: 5 neighbors

New point: $X=4.5$

All five neighbors
are labeled PASS

Predicted label = PASS




VERY simple example: 5-NN model

New point: $X=3.5$

three neighbors
are labeled PASS,
and two neighbors
are labeled FAIL

Predicted label = PASS

with probability 60%



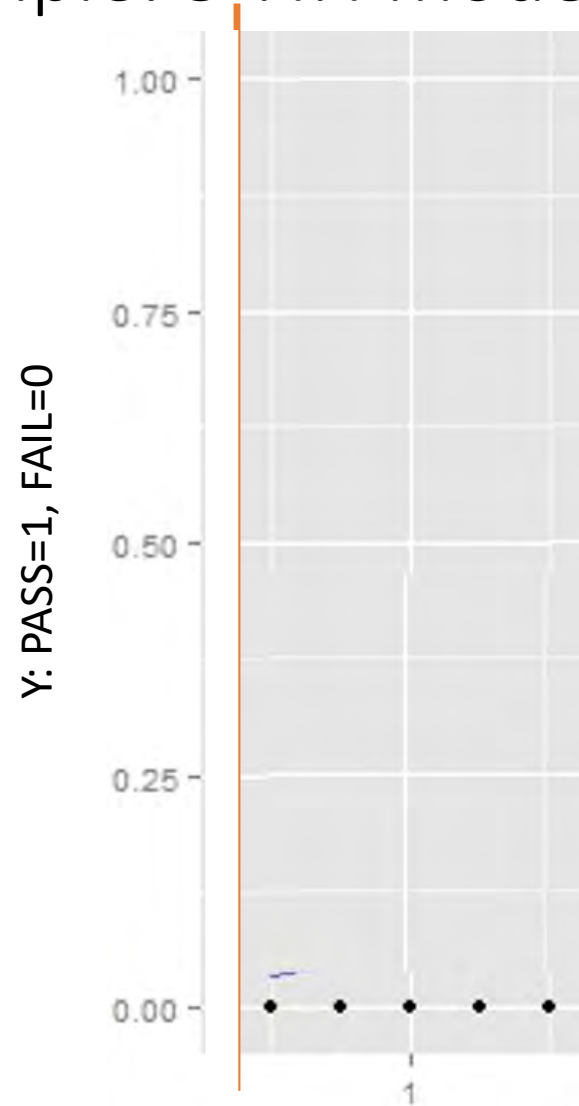
VERY simple example: 5-NN model

New point: $X=1$

All five neighbors
are labeled FAIL

Predicted label = PASS

with probability 0%

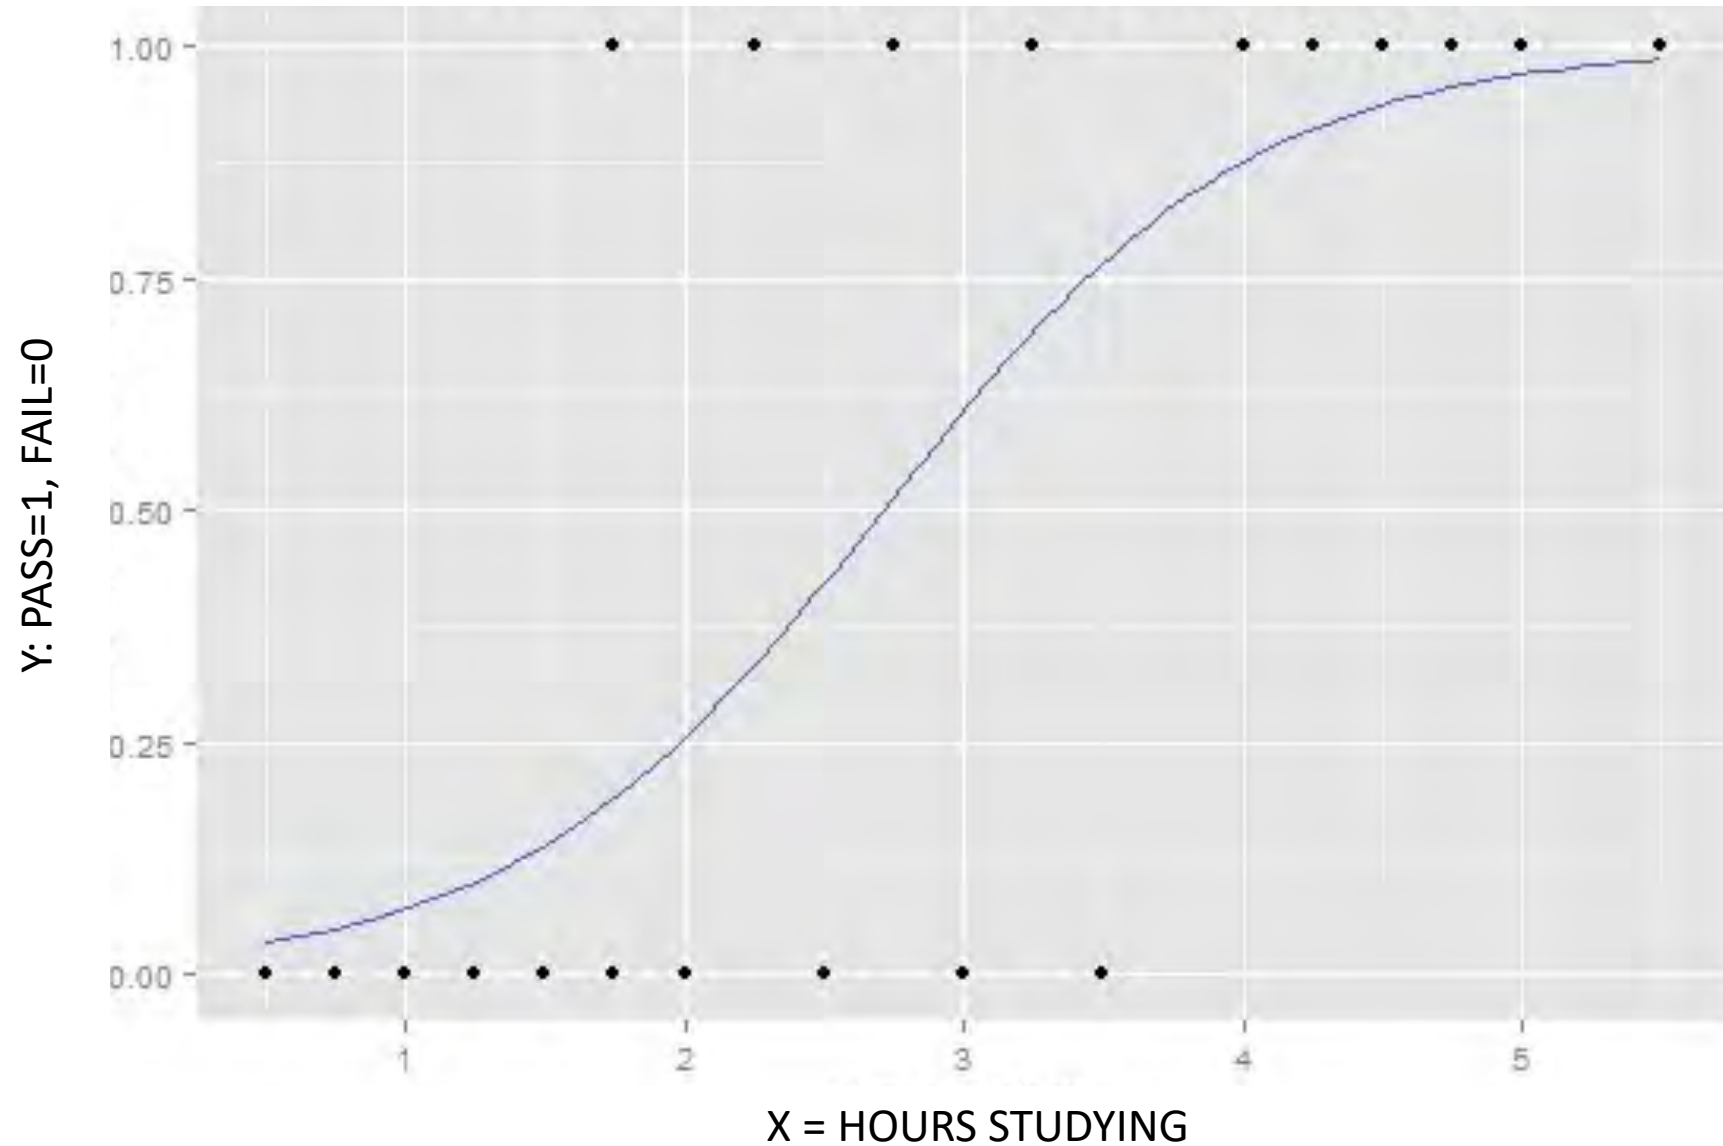


X = HOURS STUDYING

VERY simple example

Blue line is

The probability of PASS
being the correct label
for a new point X



$\text{Prob}(y)$ = Proportion of “success” among neighbors

$$\text{Prob}(y) = \frac{\sum y_i}{n} = \frac{\# \text{ of } 1\text{'s}}{\# \text{ of trials}} = \text{Proportion of "success"}$$


Goal of logistic regression: Predict the “true” proportion of success $\text{prob}(y)$ at any value of the predictor variable X

Approximated by maximizing conditional log-likelihood: $\sum \log \text{prob}(y|x)$

Model for Prob(y)

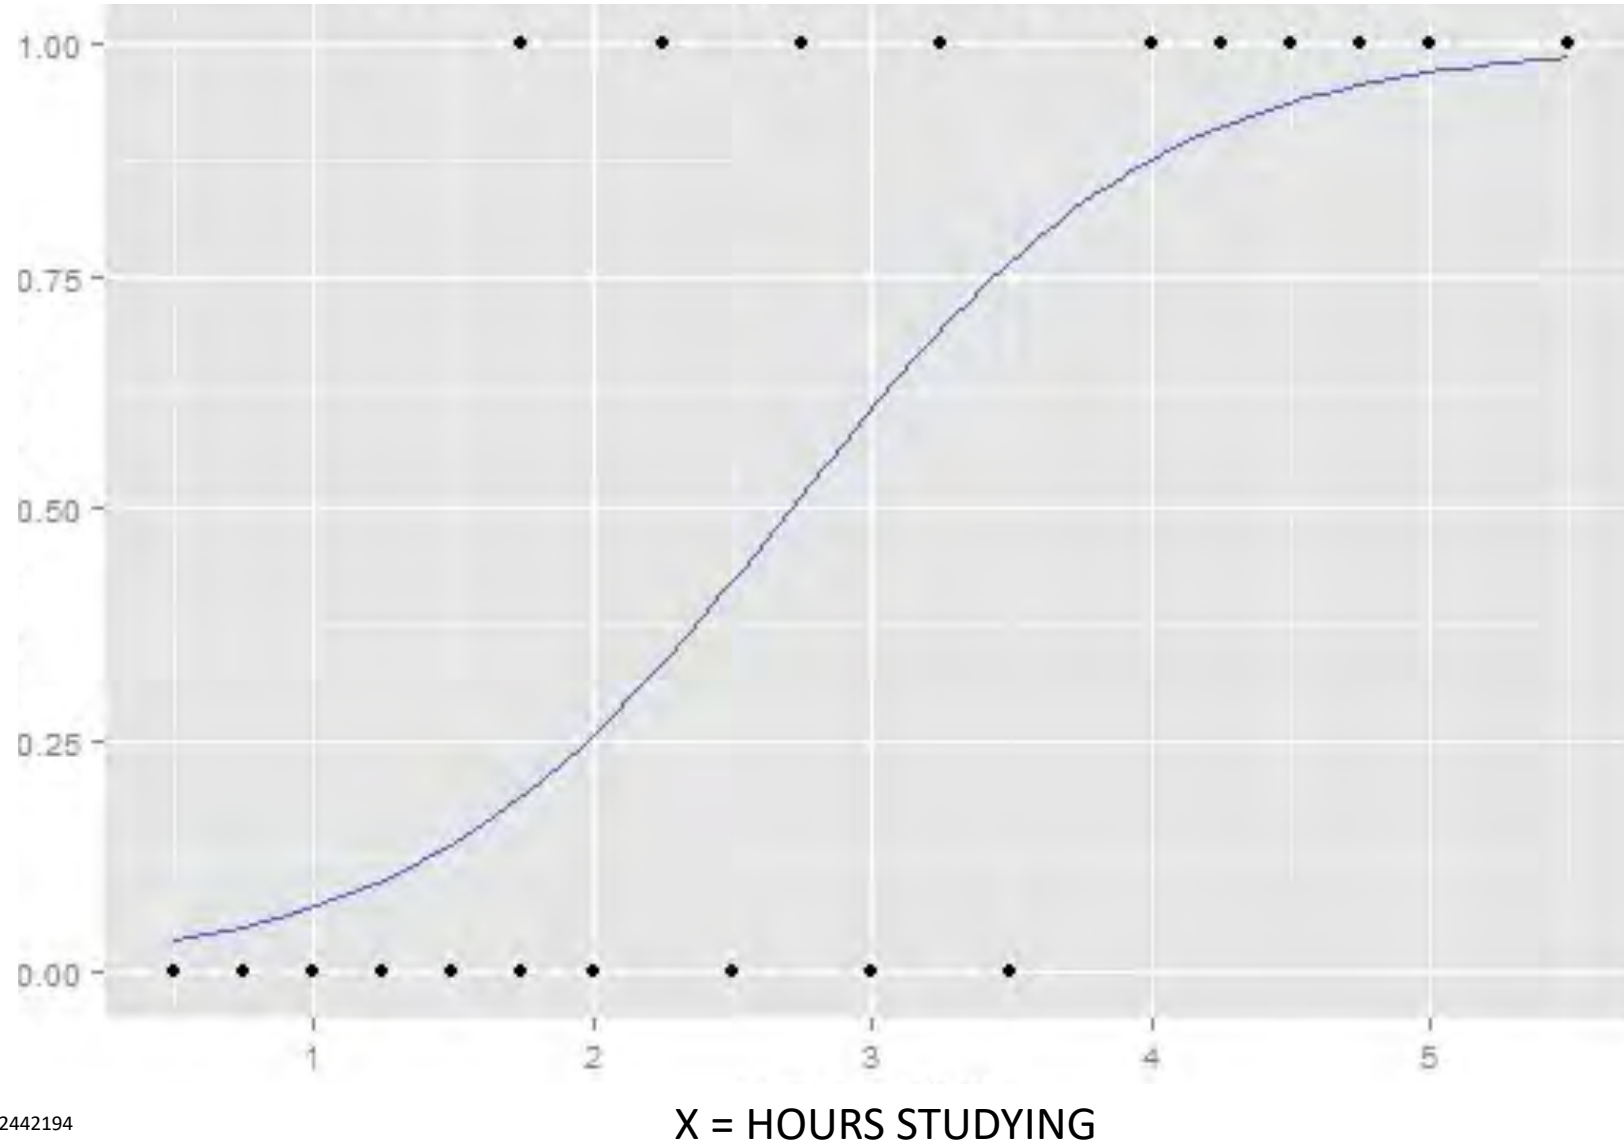
The logistic function

$$\text{Probability}(y) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$



VERY simple example

$$\text{Probability(PASS)} = \frac{1}{1 + \exp(-(1.5046 \cdot \text{Hours} - 4.0777))}$$



In practical terms

HOURS STUDYING	PASSED EXAM
4	YES
1	NO
3.5	NO
2.25	YES
0.25	NO

Known
input

Known
labels

Replace



with

HOURS STUDYING	PROB. PASS
4	%
1	%
3.5	%
2.25	%
0.25	%

Compute formula
based on logistic
function

The logistic function in more dimensions

$$\textit{Probability}(Y) = \frac{e^u}{1 + e^u} = \frac{1}{1 + e^{-u}}$$

Where u is the regular linear regression equation on M variables:


$$u = A + B_1X_1 + B_2X_2 + \dots + B_MX_M$$

Logistic Regression

Form of regression that allows the prediction of discrete variables by a mix of continuous and discrete predictors

Addresses the same questions that multiple regression does, but requires no distributional assumptions on the predictors:

- do not have to be normally distributed
- do not have to be linearly related
- do not have to have equal variance in each group



Hands-on
Example:

Logistic
regression

Logistic Regression Classifier

```
LogisticRegression(penalty='l2',  
dual=False,  
tol=0.0001,  
C=1.0,  
fit_intercept=True,  
intercept_scaling=1,  
class_weight=None,  
random_state=None,  
solver='lbfgs',  
max_iter=100,  
multi_class='auto',  
verbose=0,  
warm_start=False,  
n_jobs=None,  
l1_ratio=None)
```

https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html

Measuring how closely the formula fits the data

- **Penalty** {'l1', 'l2', 'elasticnet', 'none'}, default='l2' Used to specify the norm used in the penalization. If 'none' (not supported by the liblinear solver), no regularization is applied.
- **C** float, default=1.0 Inverse of regularization strength; must be a positive float. Smaller values specify stronger regularization.
- **l1_ratio** float, default=None The Elastic-Net mixing parameter, with $0 \leq \text{l1_ratio} \leq 1$. Only used if penalty='elasticnet'. Setting l1_ratio=0 is equivalent to using penalty='l2', while setting l1_ratio=1 is equivalent to using penalty='l1'. For $0 < \text{l1_ratio} < 1$, the penalty is a combination of L1 and L2

Choosing a method for solving the fitting problem

- **solver**{'newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'}, default='lbfgs' Algorithm to use in the optimization problem. Not every solver choice will work with every penalty choice.
- **max_iter** int, default=100 Maximum number of iterations taken for the solvers to converge.