

HPE DSI 311

Introduction to Machine Learning

Summer 2021

Instructor: Ioannis Konstantinidis

Overview

- Metrics and Scoring:
 - Confusion Matrix
 - Error Functions
 - Regularization
- Hands-on examples
 - Classification
 - Regression



What is a
model?



Linear Multivariate Regression

Statistics:

$$y = a + b_0X_0 + b_1X_1 + b_2X_2 \quad (\text{equation notation})$$

Linear Multivariate Regression

Statistics:

$$y = a + b_0 X_0 + b_1 X_1 + b_2 X_2 \quad (\text{equation notation})$$

Math:

$$\mathbf{y} = \mathbf{X} \boldsymbol{\beta} + \alpha \quad (\text{matrix notation})$$

Linear Multivariate Regression

Statistics:

$$y = a + b_0 X_0 + b_1 X_1 + b_2 X_2 \quad (\text{equation notation})$$

Math:

$$\mathbf{y} = \mathbf{X} \boldsymbol{\beta} + \alpha \quad (\text{matrix notation})$$

Computer Science:

$$y = \text{Model}(X, b, a) \quad (\text{functional notation})$$

Linear Multivariate Regression

Statistics:

$$y = a + b_0 X_0 + b_1 X_1 + b_2 X_2 \quad (\text{equation notation})$$

Math:

$$\mathbf{y} = \mathbf{X} \boldsymbol{\beta} + \alpha \quad (\text{matrix notation})$$

Computer Science:

```
Model.fit(X,y)           (object oriented notation)  
y = Model.predict(X)
```


Linear Multivariate Regression

Statistics:

$$y = a + b_0 X_0 + b_1 X_1 + b_2 X_2 \quad (\text{equation notation})$$


Math:

$$\mathbf{y} = \mathbf{X} \boldsymbol{\beta} + \alpha \quad (\text{matrix notation})$$

Computer Science:

```
Model.fit(X,y)           (object oriented notation)  
y = Model.predict(X)
```

Deep Learning/
ML:



(network notation)

You say to-may-to, I say to-mah-to

Math:

y, X -> variables

α, β -> parameters

Statistics:

y, X_i -> variables

a, b_i -> parameters

Computer Science:

X, y

-> parameters when fitting

b, a

-> parameters when predicting

called weights w for networks

How do you
“fit” a
model?



Assessing model fitness – supervised ML

Calculate the **parameter values** that make model predictions fit the training data **most closely**

Assessing model fitness – supervised ML

Calculate the **parameter values** that make model predictions match the training data **most closely**

Naïve solution: Exhaustive Search

```
Input: X_train, y_train, Model
```

```
For [b, a] in someParameterSpace
```

```
    y_predict <- Model( X_train, b, a )
```

```
    penaltyList.append <- Penalty( y_predict, y_train )
```

```
Output: [b_fit, a_fit] = argmin(penaltyList)
```

Assessing model fitness – supervised ML

Calculate the **parameter values** that make the model predictions match the training data **most closely**

Naïve solution: Exhaustive Search

```
Input: X_train, y_train, Model

For [b, a] in someParameterSpace
    y_predict          <- Model( X_train, b, a )
    penaltyList.append <- Penalty( y_predict, y_train )

Output: [b_fit, a_fit] = argmin(penaltyList)
```

Need an appropriate Penalty() for y

Assessing model fitness – unsupervised ML

Calculate the training **data points** that **are closest to** the new data point

Assessing model fitness – unsupervised ML

Calculate the training **data point** that **is closest to** the new data point

Need an appropriate `Penalty()` for X

Assessing model fitness

Penalty() for y : compare N pairs (y_{predict} , y_{train})

- Metric / Objective / Cost / Loss function

Penalty() for X : compare two N -dimensional points

- Similarity / Affinity

Testing model performance

Penalty() for y : compare N pairs (y_{predict} , y_{train})

- Scoring / Error function

Penalty() for X : compare two N -dimensional points

- Distance

Penalty()

Metric function for assessing fitness during training

Scoring function for testing performance during
evaluation

How to define penalty functions



Predictive Capability for Classification Tasks

accepting (word
article).

focus n point


converging rays of light,
heat, wave of sound, magnet.

centre of activity or
intensity; pl foci; v
adjust; cause to converge;
concentrate; a focal
pertaining to focus

Confusion Matrix



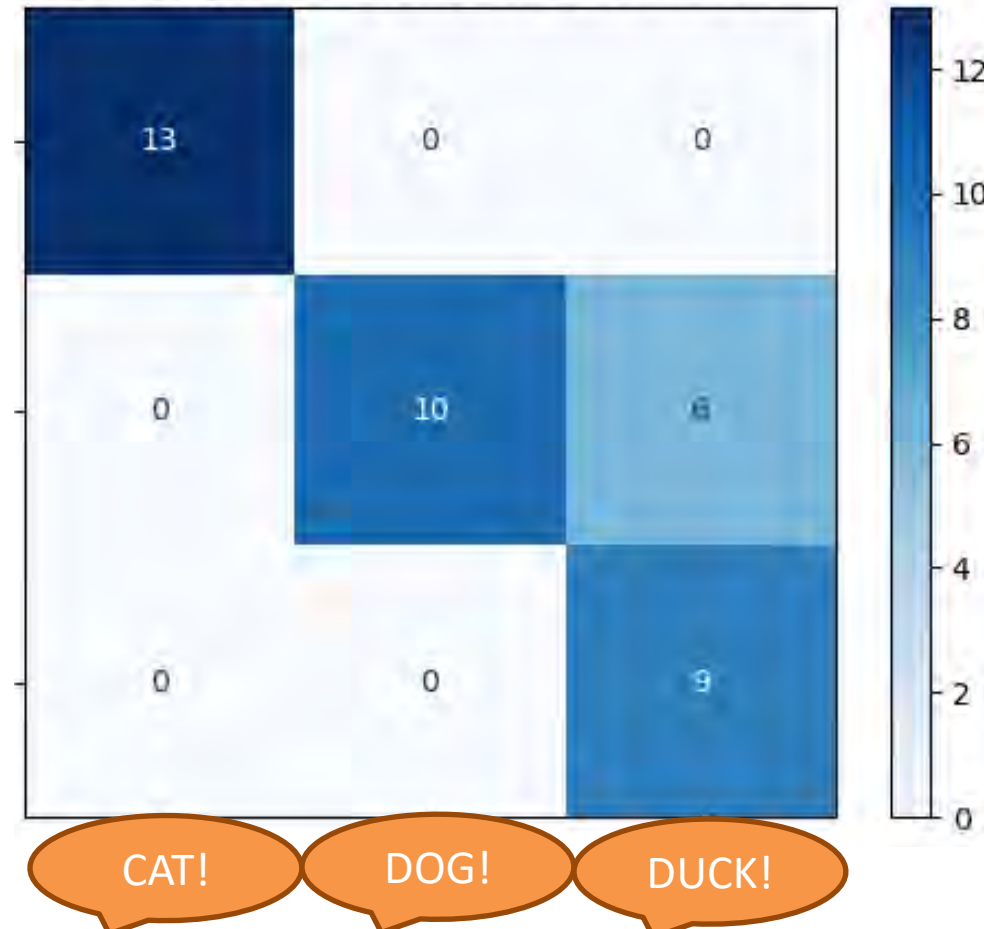

Confusion Matrix



Confusion Matrix

Columns: predictions made by the classifier (labels y)

Rows: actual observations (points X)



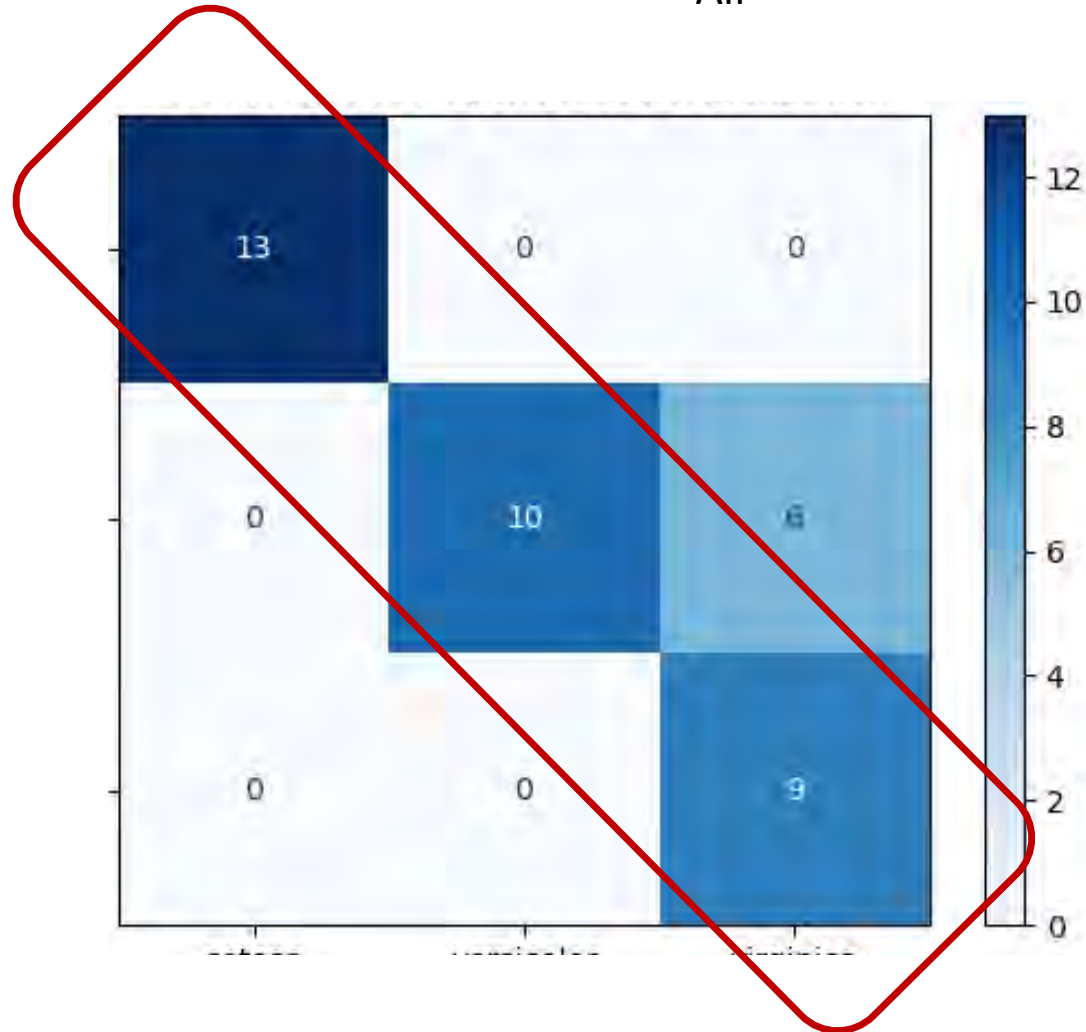
Confusion Matrix

- Diagonal: # of points for which predicted label = true label
- Off-diagonal: # of points that are mislabeled by the classifier
- The higher the diagonal values of the confusion matrix, the better



Confusion Matrix


$$\text{Accuracy} = \frac{\text{Diagonal}}{\text{All}}$$



Focus on a single label: is it a cat?

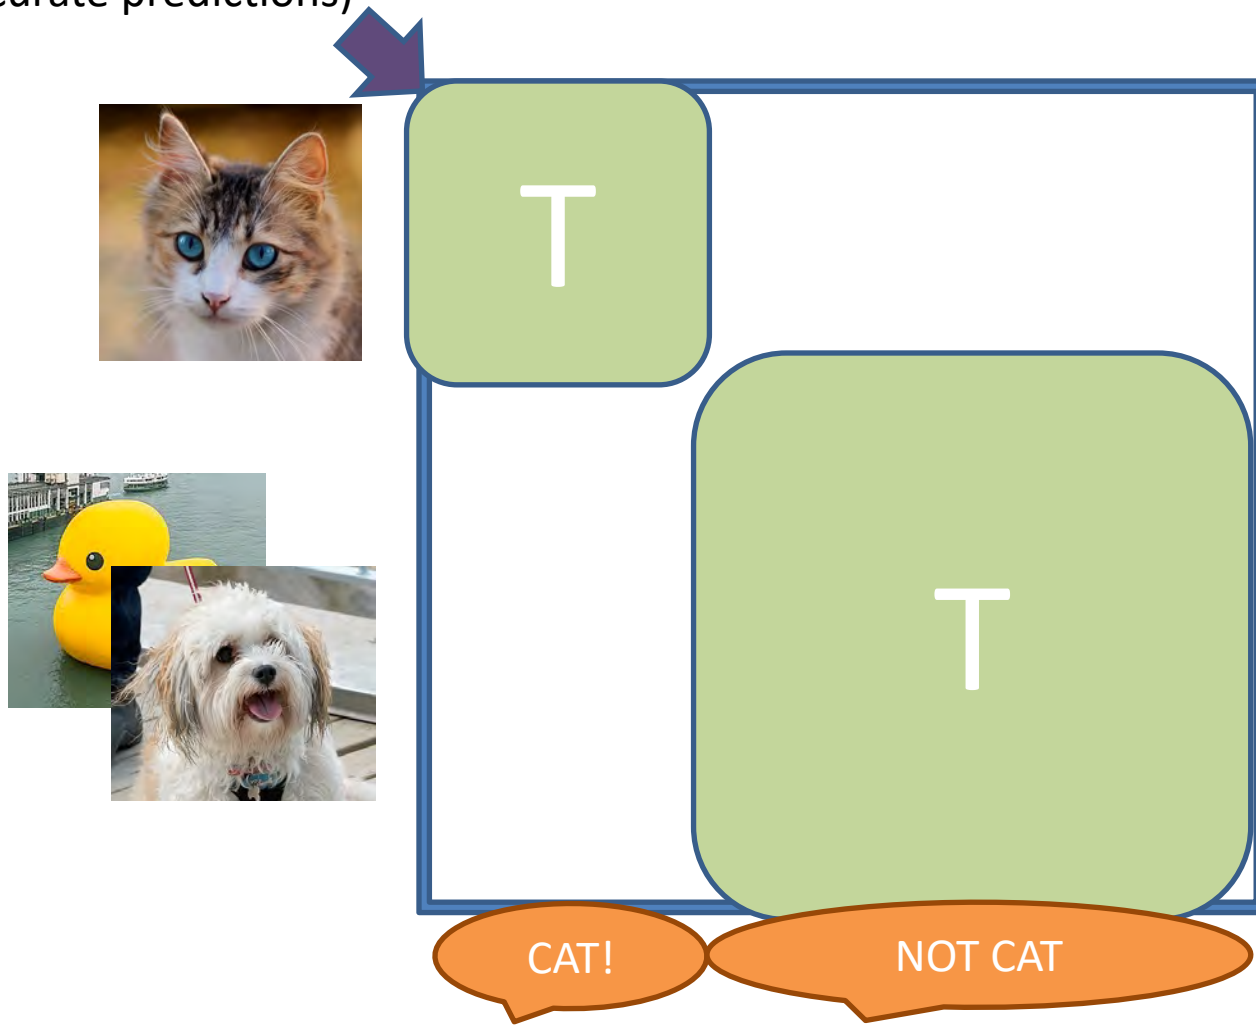
Positives column
(predicted to
be the target)

Negatives column
(predicted NOT to
be the target)

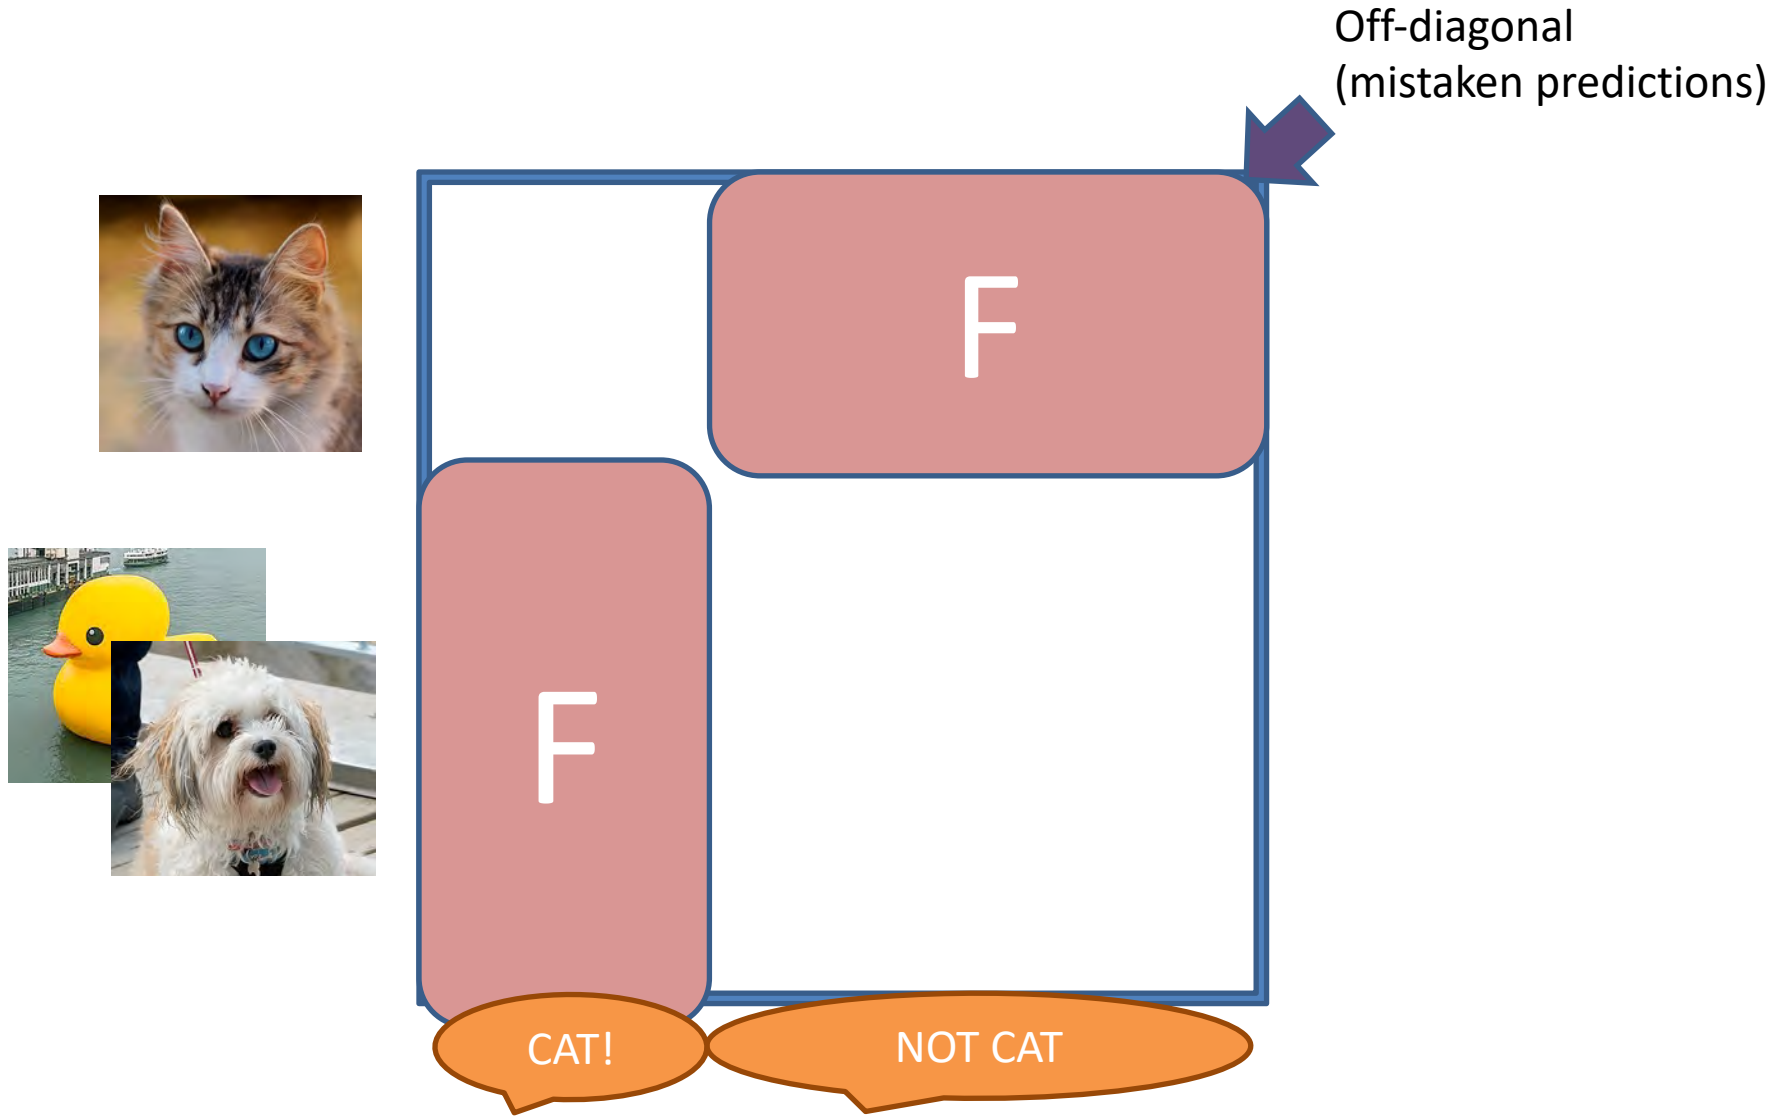


Focus on a single label: is it a cat?

Diagonal
(accurate predictions)



Focus on a single label: is it a cat?



Focus on a single label: is it a cat?

Columns: predictions made by the classifier (labels y)

Rows: actual observations (points X)


T+	F-
F+	T-

Focus on a single label: is it a cat?




Focus on a single label: Precision


$$\text{Precision} = \frac{\boxed{T+}}{\boxed{+}}$$



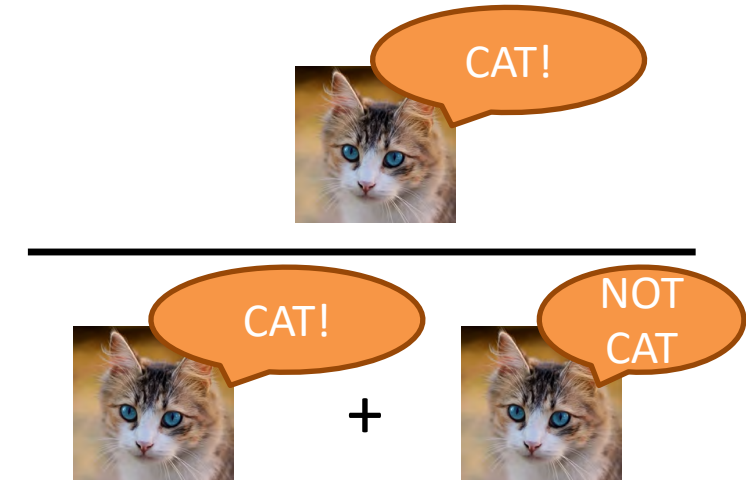
Positives column
(predicted to
be the target)



Focus on a single label: Recall

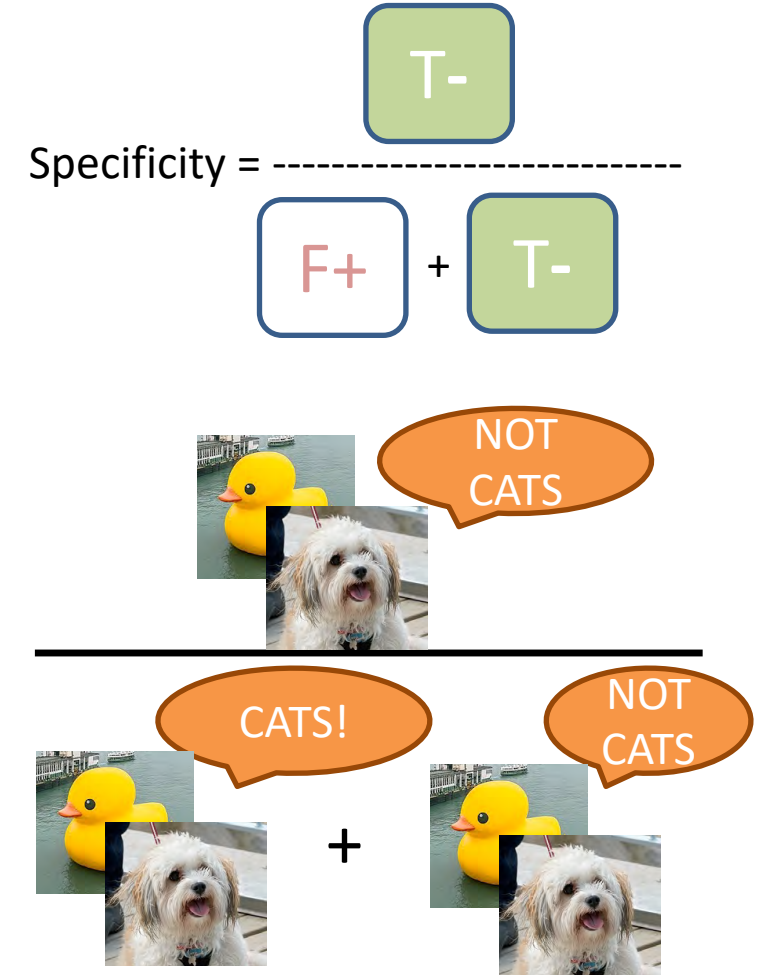



$$\text{Recall} = \frac{\text{T+}}{\text{T+} + \text{F-}}$$



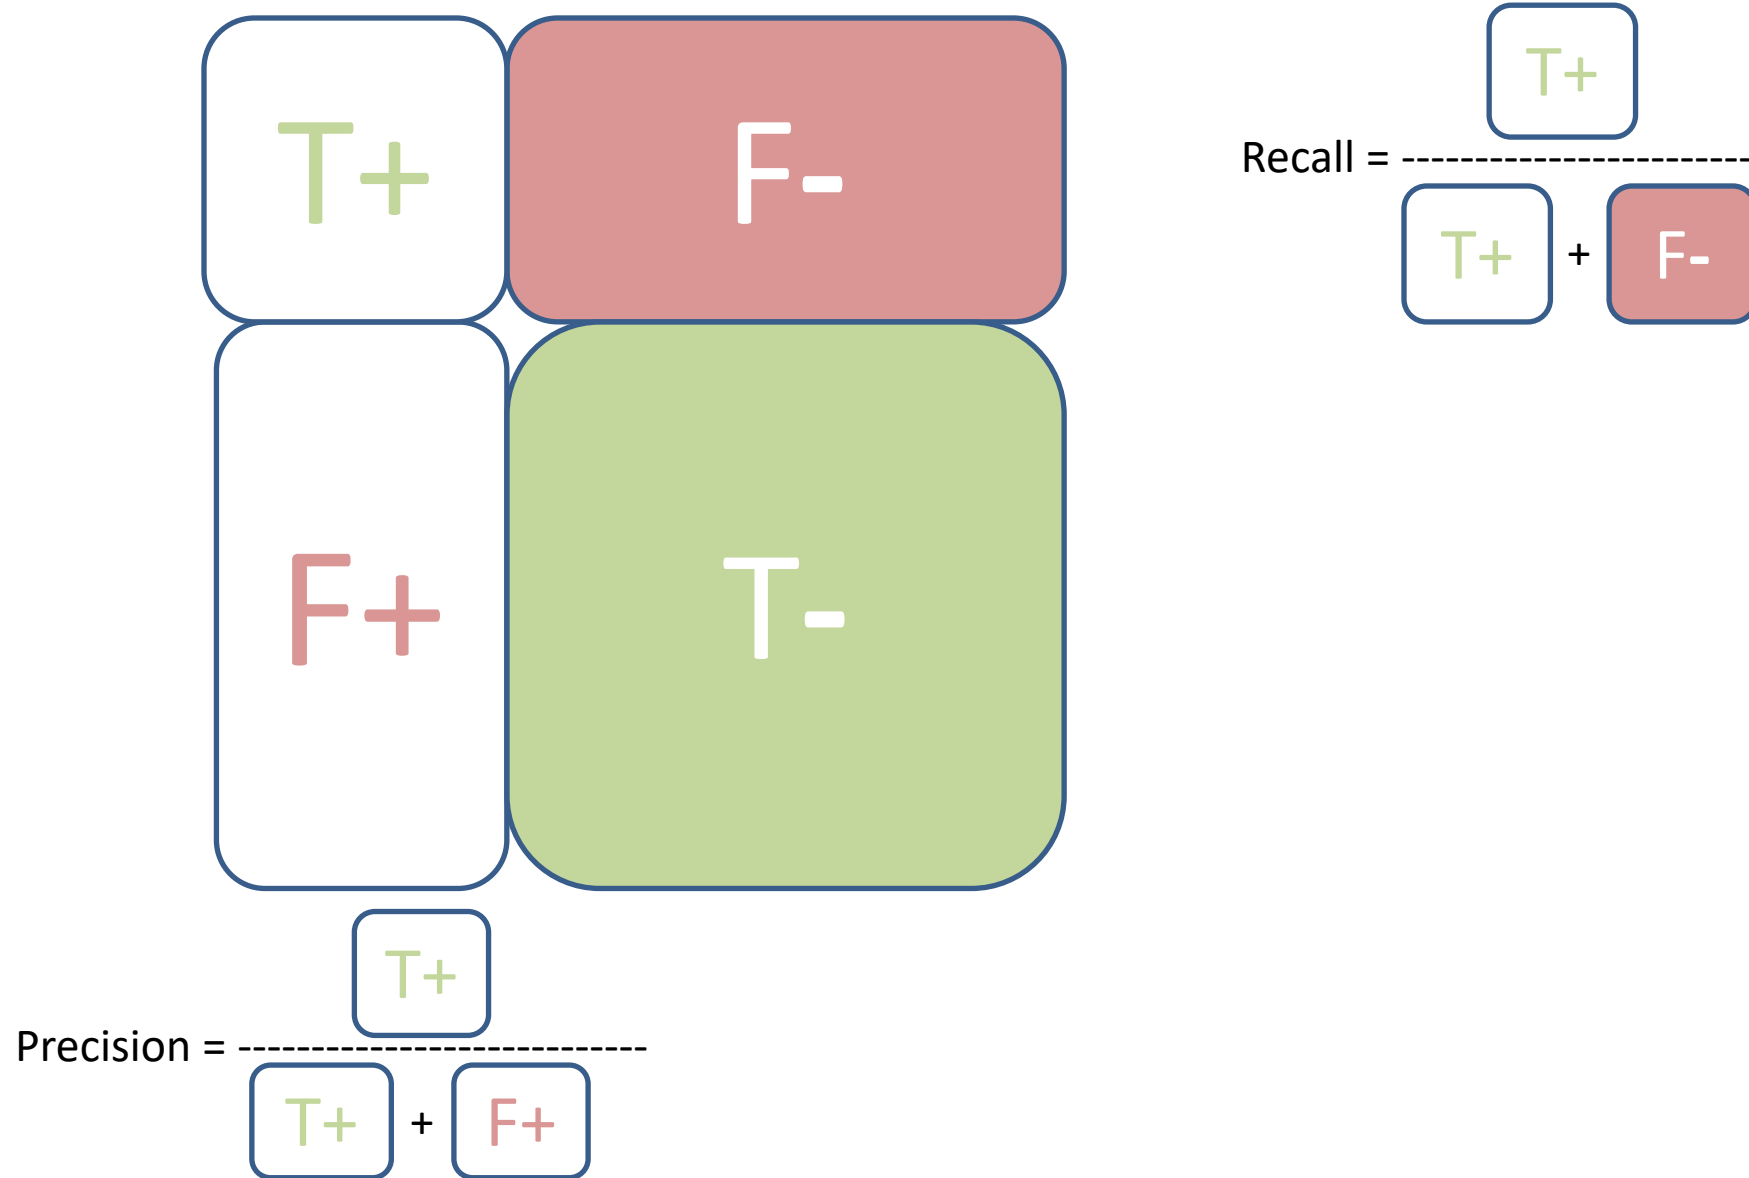
AKA sensitivity, hit rate

Focus on a single label: specificity




AKA selectivity

Focus on a single label: combinations

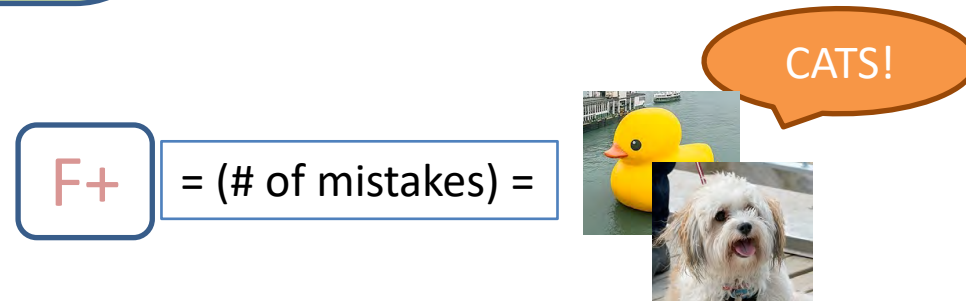



Focus on a single label: combinations




$$\text{Precision} = \frac{\text{T+}}{\text{T+} + \text{F+}}$$

$$\text{Recall} = \frac{\text{T+}}{\text{T+} + \text{F-}}$$



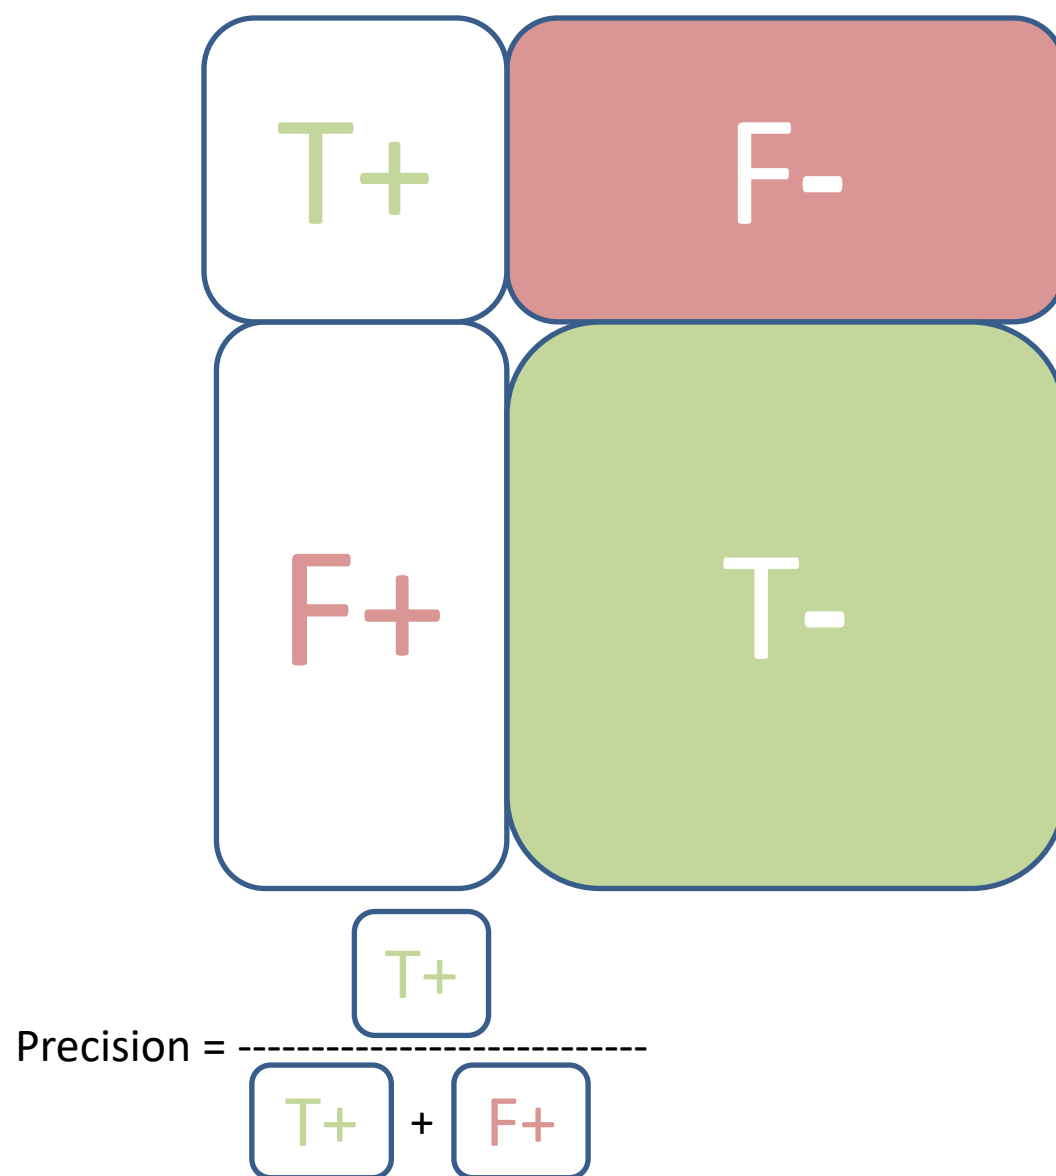
Focus on a single label: combinations



$$\text{Recall} = \frac{\text{T+}}{\text{T+} + \text{\# of misses}}$$


$$F_1 = \frac{\text{T+}}{\text{T+} + \frac{\text{\# of mistakes} + \text{\# of misses}}{2}}$$

Focus on a single label: combinations



$$\text{Recall} = \frac{\text{T+}}{\text{T+} + \text{F-}}$$

$$F_1 = \frac{1}{\frac{1}{2} \left(\frac{1}{\text{recall}} + \frac{1}{\text{precision}} \right)}$$
$$= 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



Hands-on
Example:

Classification
using k-NN +
Logistic
Regression

Confusion matrix

```
Plot_confusion_matrix(estimator, X, y_true,  
labels=None,  
sample_weight=None,  
normalize=None,  
display_labels=None,  
include_values=True,  
xticks_rotation='horizontal',  
values_format=None,  
cmap='viridis',  
ax=None)
```

https://scikit-learn.org/stable/modules/model_evaluation.html#confusion-matrix

Confusion matrix

Labels: List of labels to index the matrix. This may be used to reorder or select a subset of labels. If None is given, those that appear at least once in y_true or y_pred are used in sorted order.

Normalize: Normalizes confusion matrix over the true (rows), predicted (columns) conditions or all the population. If None, confusion matrix will not be normalized.

include_values: Includes values in confusion matrix.

Classification Report

```
classification_report(y_true, y_pred,  
labels=None,  
target_names=None,  
sample_weight=None,  
digits=2,  
output_dict=False,  
zero_division='warn')
```

https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification_report.html

Classification Report

'macro': Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

'weighted': Calculate metrics for each label, and find their average weighted by support (the number of true instances for each label). This alters 'macro' to account for label imbalance; it can result in an F-score that is not between precision and recall.

Note that if all labels are included, “micro”-averaging in a multiclass setting will produce precision and recall scores that are all identical to accuracy.

A magnifying glass is held over an open dictionary. The lens is focused on the word 'focus', which is highlighted with a green marker. The text around the word is slightly blurred, but some words like 'accepting', 'article', 'focus n point', 'converging rays of light', 'heat, wave of sound, light', 'centre of activity or intensity', 'adjust; cause to converge', 'concentrate; a focal', and 'pertaining to focus' are visible. The title 'Predictive Capability for Regression Tasks' is overlaid in white text on the magnifying glass.

Predictive Capability for Regression Tasks

Mean Squared Error (MSE)

$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

where N is the number of data points,
 f_i the value returned by the model and
 y_i the actual value for data point i .

Mean Squared Error (MSE)

$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$


where N is the number of data points,
 f_i the value returned by the model and
 y_i the actual value for data point i .

Euclidean distance squared,
divided by number of points

Mean Squared Error (MSE)

$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

where N is the number of data points,
 f_i the value returned by the model and
 y_i the actual value for data point i .



Mean Squared Error (MSE)

$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

where N is the number of data points,
 f_i the value returned by the model and
 y_i the actual value for data point i .

N=2

2	2.236	2.828
1	1.414	2.236
(y ₁ , y ₂)	1	2

Mean Absolute Deviation (MAD)

$$\frac{1}{N} \sum_{i=1}^N |f_i - y_i|$$

Mean Absolute Deviation (MAD)

Manhattan distance
divided by number of points

$$\frac{1}{N} \sum_{i=1}^N |f_i - y_i|$$

N=2

2	3	4
1	2	3
(y ₁ , y ₂)	1	2

Maximum error

N=2

	2	2	2
	1	1	2
(y_1, y_2)	1	2	2

Recall the L^p norms

$$\|x\|_p = (|x_1|^p + |x_2|^p + \cdots + |x_n|^p)^{1/p}$$

$$\|x\|_\infty = \max\{|x_1|, |x_2|, \dots, |x_n|\}$$



Unit circle for different values of p

Regularization: mix and match

accepting (word
article).

focus n point

converging rays of light,
heat, waves, etc. meet;

centre of activity or
intensity; p. focuses, focus v
adjust; cause to converge;
concentrate; a focal
pertaining to focus

Mix and match

Multivariate Regression: $F = X \beta + \text{constant}$

$$F = A + B_1 X_1 + B_2 X_2 + \dots + B_K X_K$$

Mix and match

Multivariate Regression: $F = X \beta + \text{constant}$

$$\sum_{i=1}^N (f_i - y_i)^2$$

$$= (y - X\beta)^T (y - X\beta)$$

Mix and match

Multivariate Regression: $F = X \beta + \text{constant}$

$$\text{Ridge Cost} = (y - X\beta)^T (y - X\beta) + \alpha ||\beta||_2^2$$


$$\text{Lasso Cost} = (y - X\beta)^T (y - X\beta) + \alpha ||\beta||_1$$

α is the
regularization
(hyper)parameter

Mix and match

Multivariate Regression: $F = X \beta + \text{constant}$

	L^2		L^2
Ridge Cost =	$(y - X\beta)^T (y - X\beta)$	+	$\alpha \beta _2^2$
Lasso Cost =	$(y - X\beta)^T (y - X\beta)$	+	$\alpha \beta _1$
	L^2		L^1



Hands-on
Example:

Linear
Regression